

Deep Learning for Anomaly Detection



Meryem Janati Idrissi

Ismail Berrada

Table of Contents



01

...

Deep Learning & Anomaly Detection



02

...

Deep Structures for AD

2-1- Generative Adversarial Networks

2-2- Autoencoders

2-3- Variational Autoencoders

2-4- Adversarial Autoencoders

2-5- Normalizing Flows



03

...

Evaluation

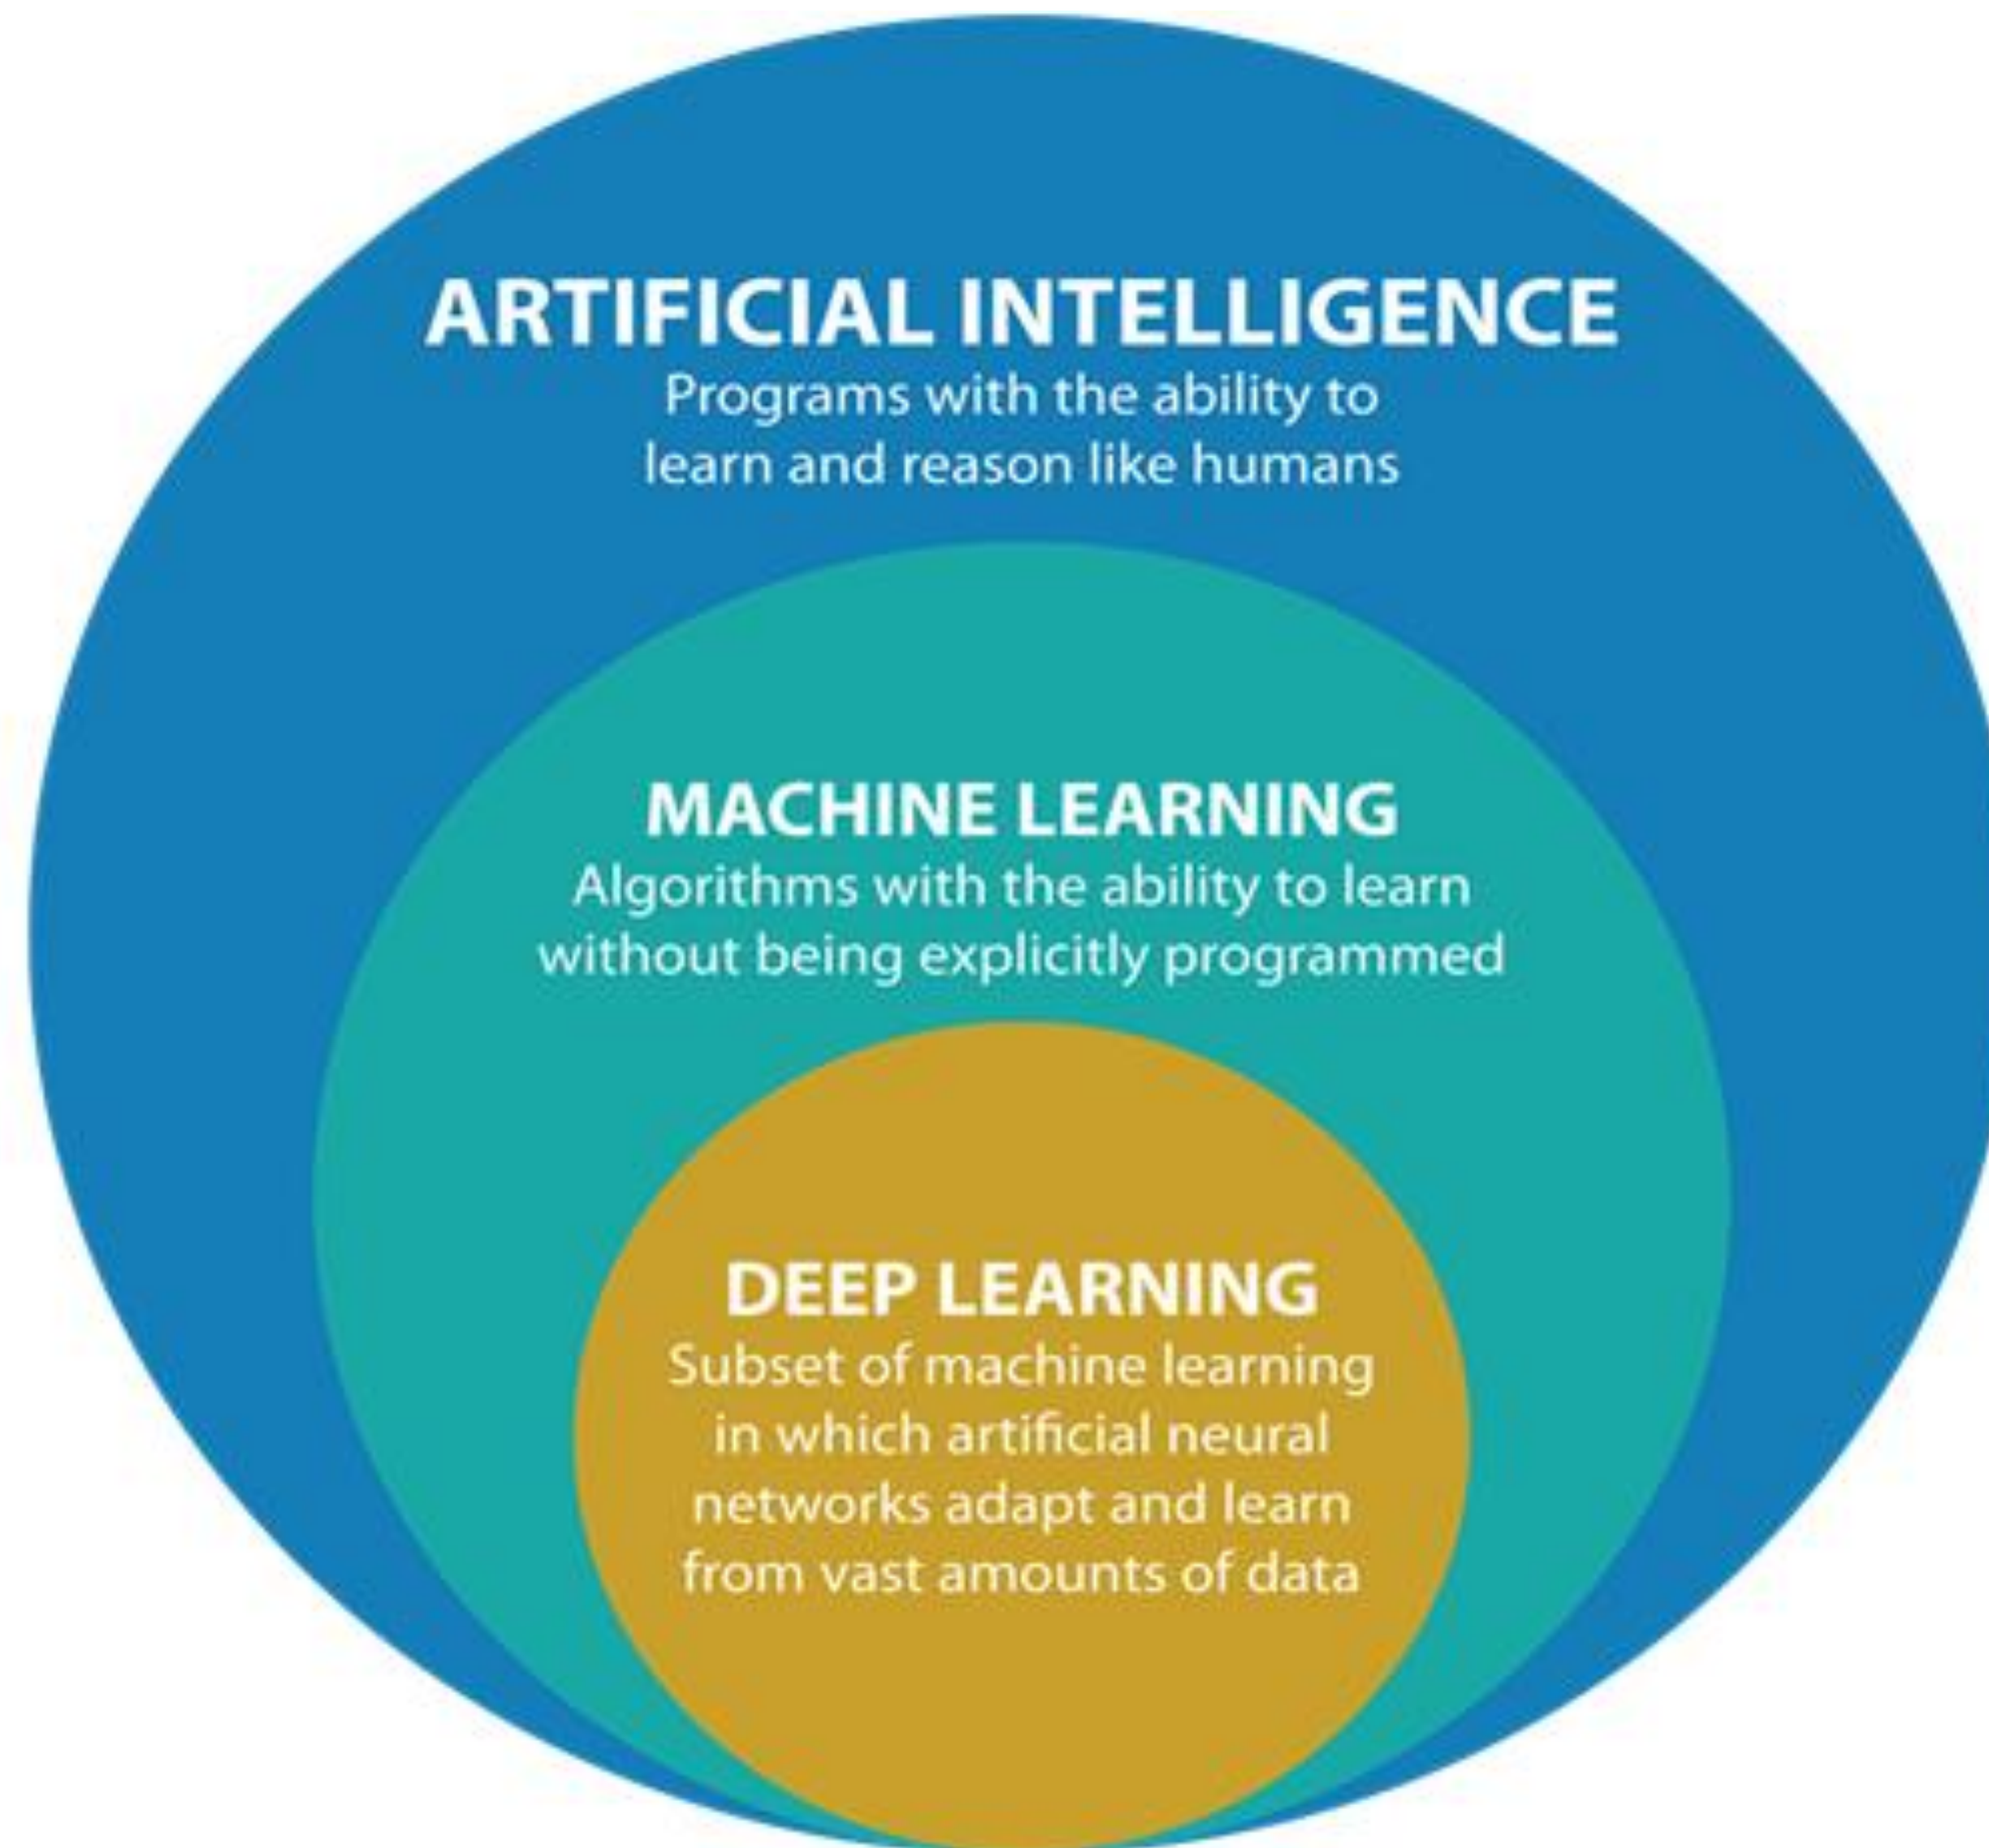


04

...


Conclusion & Perspectives

Deep Learning




Anomaly Detection

- In recent years, anomaly detection has become increasingly important in a variety of domains in business, science and technology.
- Anomalous events occur relatively infrequently, however, when they do occur, their consequences can be quite dramatic.



What is an Anomaly?

- Anomalies are data points within the datasets that appear to deviate markedly from expected outputs.
- An anomaly is an observation which deviates so much from the other observations which raises suspicion of it being generated through a different mechanism.
- Anomaly Detection is to identify patterns in a given dataset that don't confirm to expected behavior.



Applications of AD

Applications of anomaly detection include fraud detection in financial transactions, fault detection in manufacturing, intrusion detection in a computer network, monitoring sensor readings in an aircraft, spotting potential risk or medical problems in health data, and predictive maintenance.



What makes AD a hard problem?

- Lack of labelled data. Anomalies are hard to detect manually. As a result , labeling them is not feasible.
- Traditional methods don't scale as the size of data increases.
- Traditional ML methods have limitations in extracting features from the high dimensional data. A lot of manual input is required to make the data ready for machine learning.




Why Deep learning for Anomaly Detection?

- Deep learning is highly scalable, especially because of modern distributed computing methods.
- Deep Learning methods have proven capabilities in image, text and speech recognition.
- They can efficiently learn information/features from data with minimum human intervention.

Generative Adversarial Network (GAN)



A Fraud (Generator)

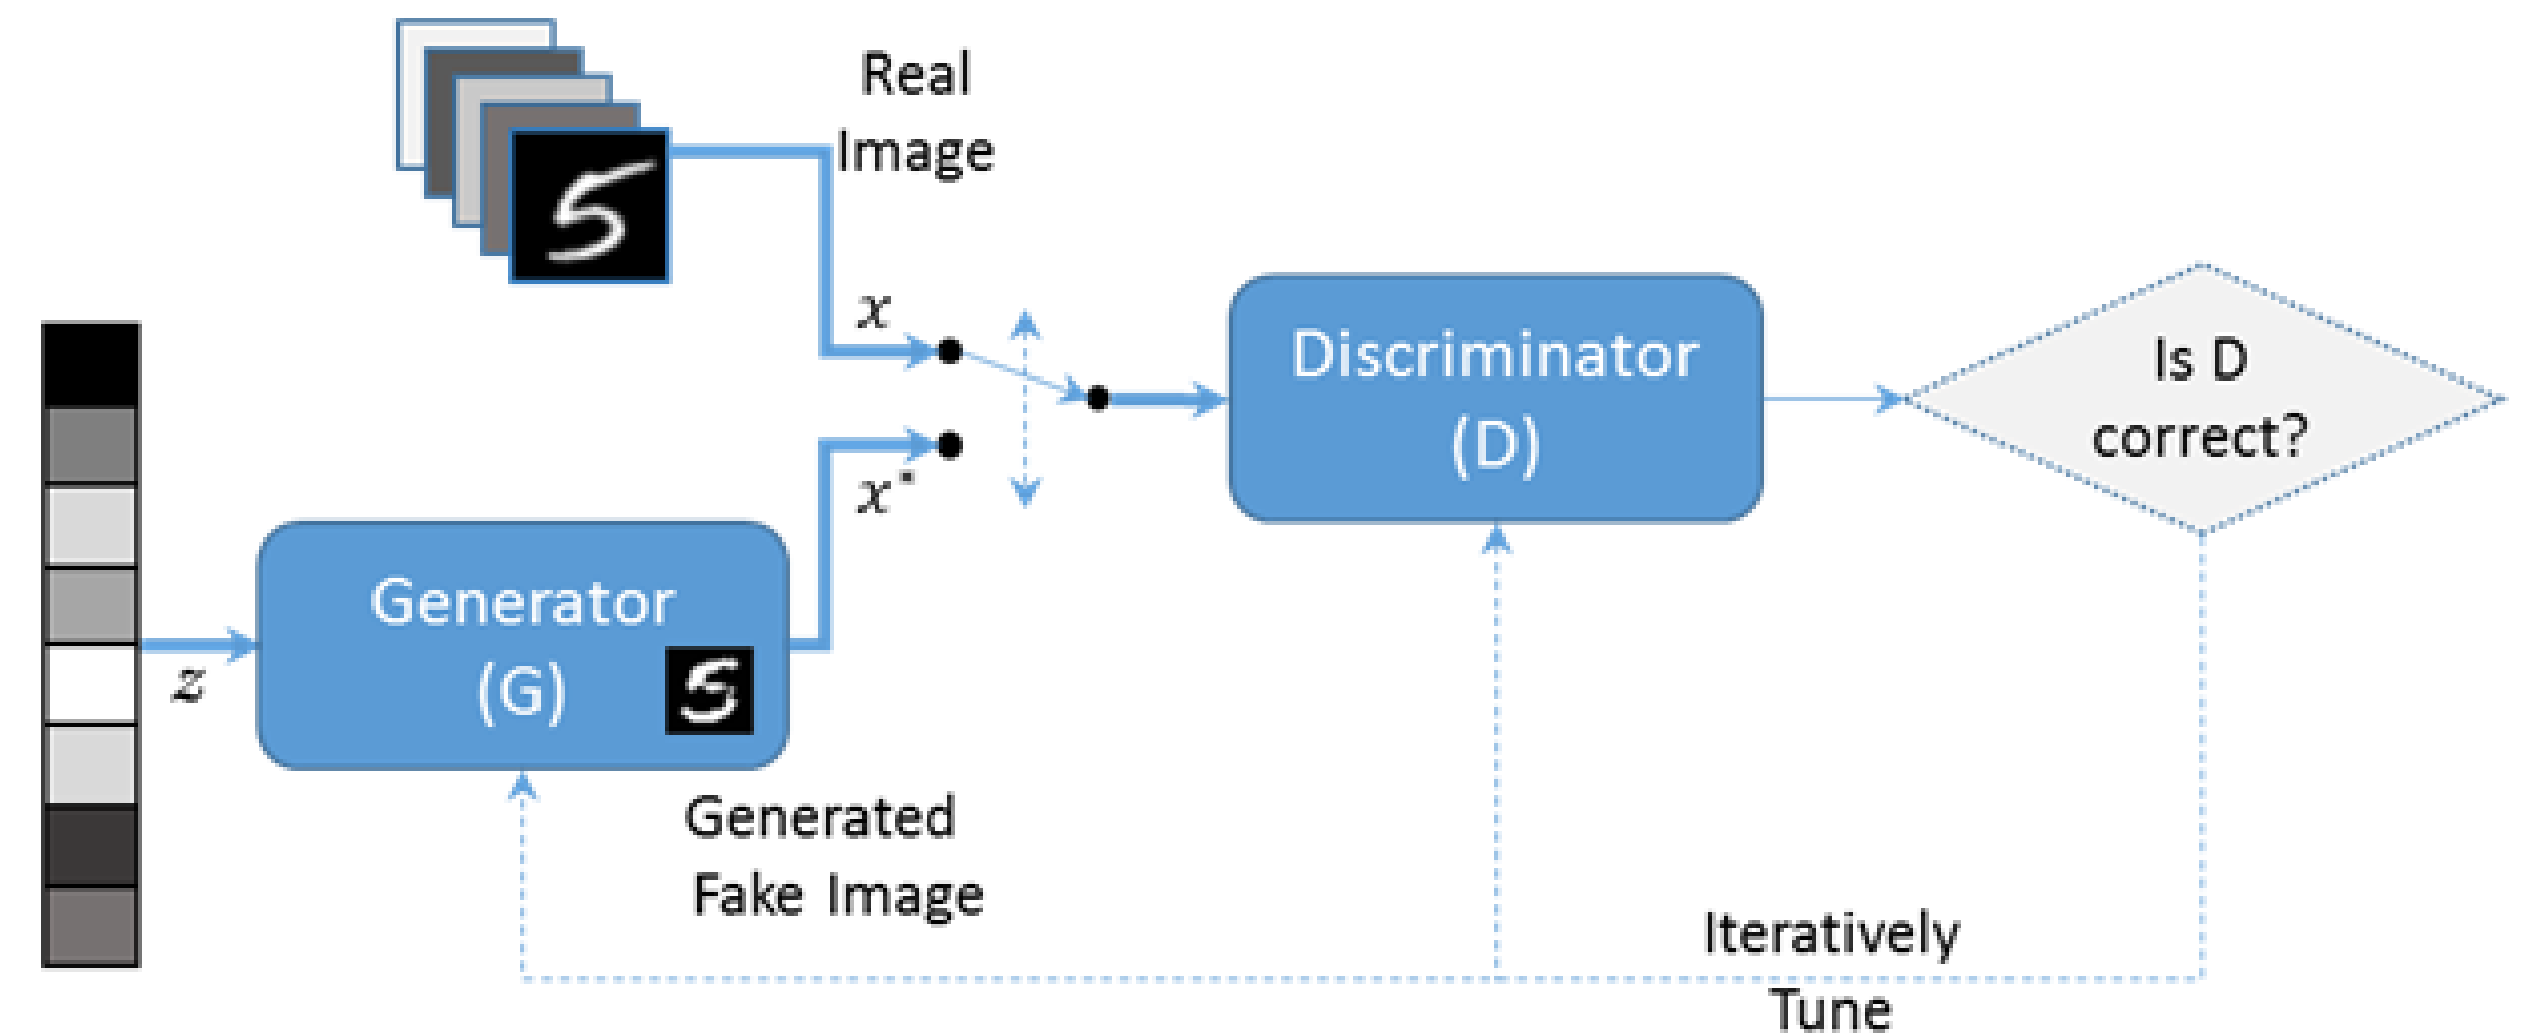


Cop (Discriminator)

How does GAN works

Two components:

- **Generator (G):** Learns input data X representation.
- **Discriminator (D):** Distinguish input data and samples from generator.



D and G play a two-player minimax game with value function $V(G,D)$.

$$\min_G \max_D V(D, G) = E_{x \sim p_{\text{data}}(x)} [\log D(x)] + E_{z \sim p_z(z)} [1 - D(G(z))]$$

Autoencoders (AE)

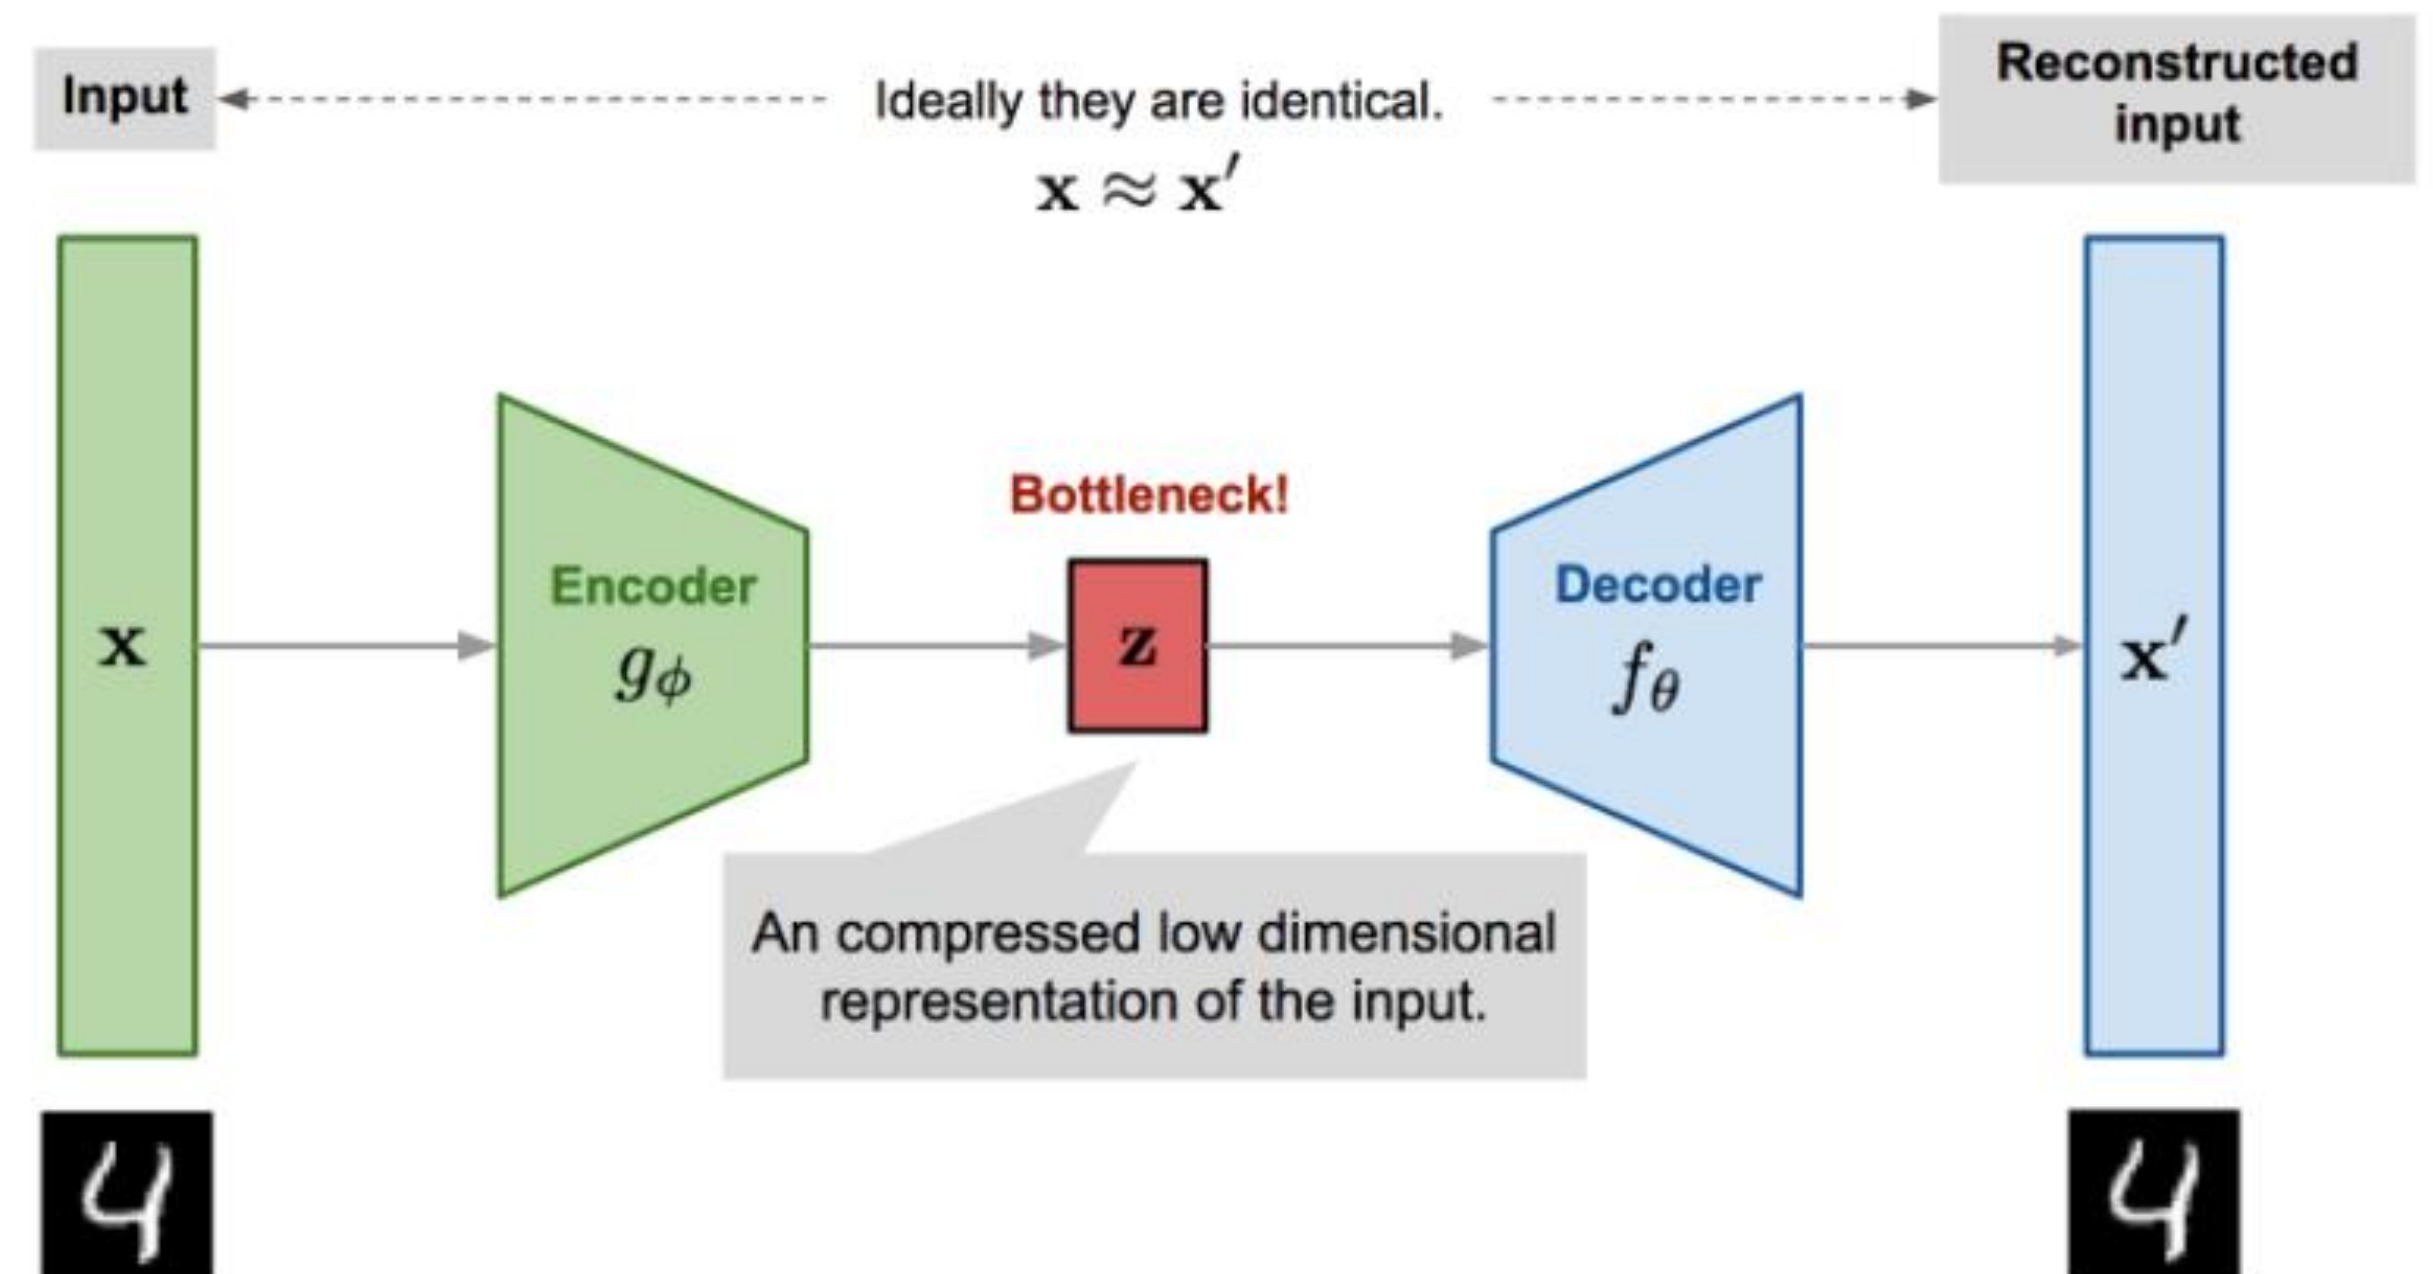
Two components:

- **Encoder:** Learns a representation or an encoding of a set of data in an unsupervised manner.
- **Decoder:** it recovers the data from the code, likely with larger and larger output layers.

The goal is to minimize the reconstruction loss $L(X, X')$, the differences between the original input and the reconstruction.

MSE loss:

$$L_{AE}(\theta, \phi) = \frac{1}{n} \sum_{i=1}^n (x^{(i)} - f_{\theta}(g_{\phi}(x^{(i)})))^2$$




Variational Autoencoders (VAE)

Two components:

- **Encoder:** Learns a representation or an encoding of a set of data in an unsupervised manner.
- **Decoder:** it recovers the data from the code, likely with larger and larger output layers.

The goal is to maximize the marginal likelihood:


$$\begin{aligned}\log p(x) &\geq \log p(x) - KL(q_\phi(z|x) || p(z|x)) \\ &= E_{z \sim q_\phi(z|x)} \log p_\theta(x|z) - KL(q_\phi(z|x) || p(z))\end{aligned}$$



Adversarial Autoencoders (AAE)

Three components:

- **Encoder(Generator):** the encoder will take the input and transform it into a lower dimension (latent code z).
- **Decoder:** the decoder will take the latent code z and transform it into the generated image.
- **Discriminator:** the discriminator takes random vector z sampled from the chosen distribution (real) and also the encoded latent code z (fake) from the autoencoder as the input. It will check whether the input is real or not.



The goal :

$$\min_E \max_D V(E, D) = E_{z \sim p(z)} [\log D(z)] + E_{x \sim p_{\text{data}}(x)} [1 - D(E(x))]$$

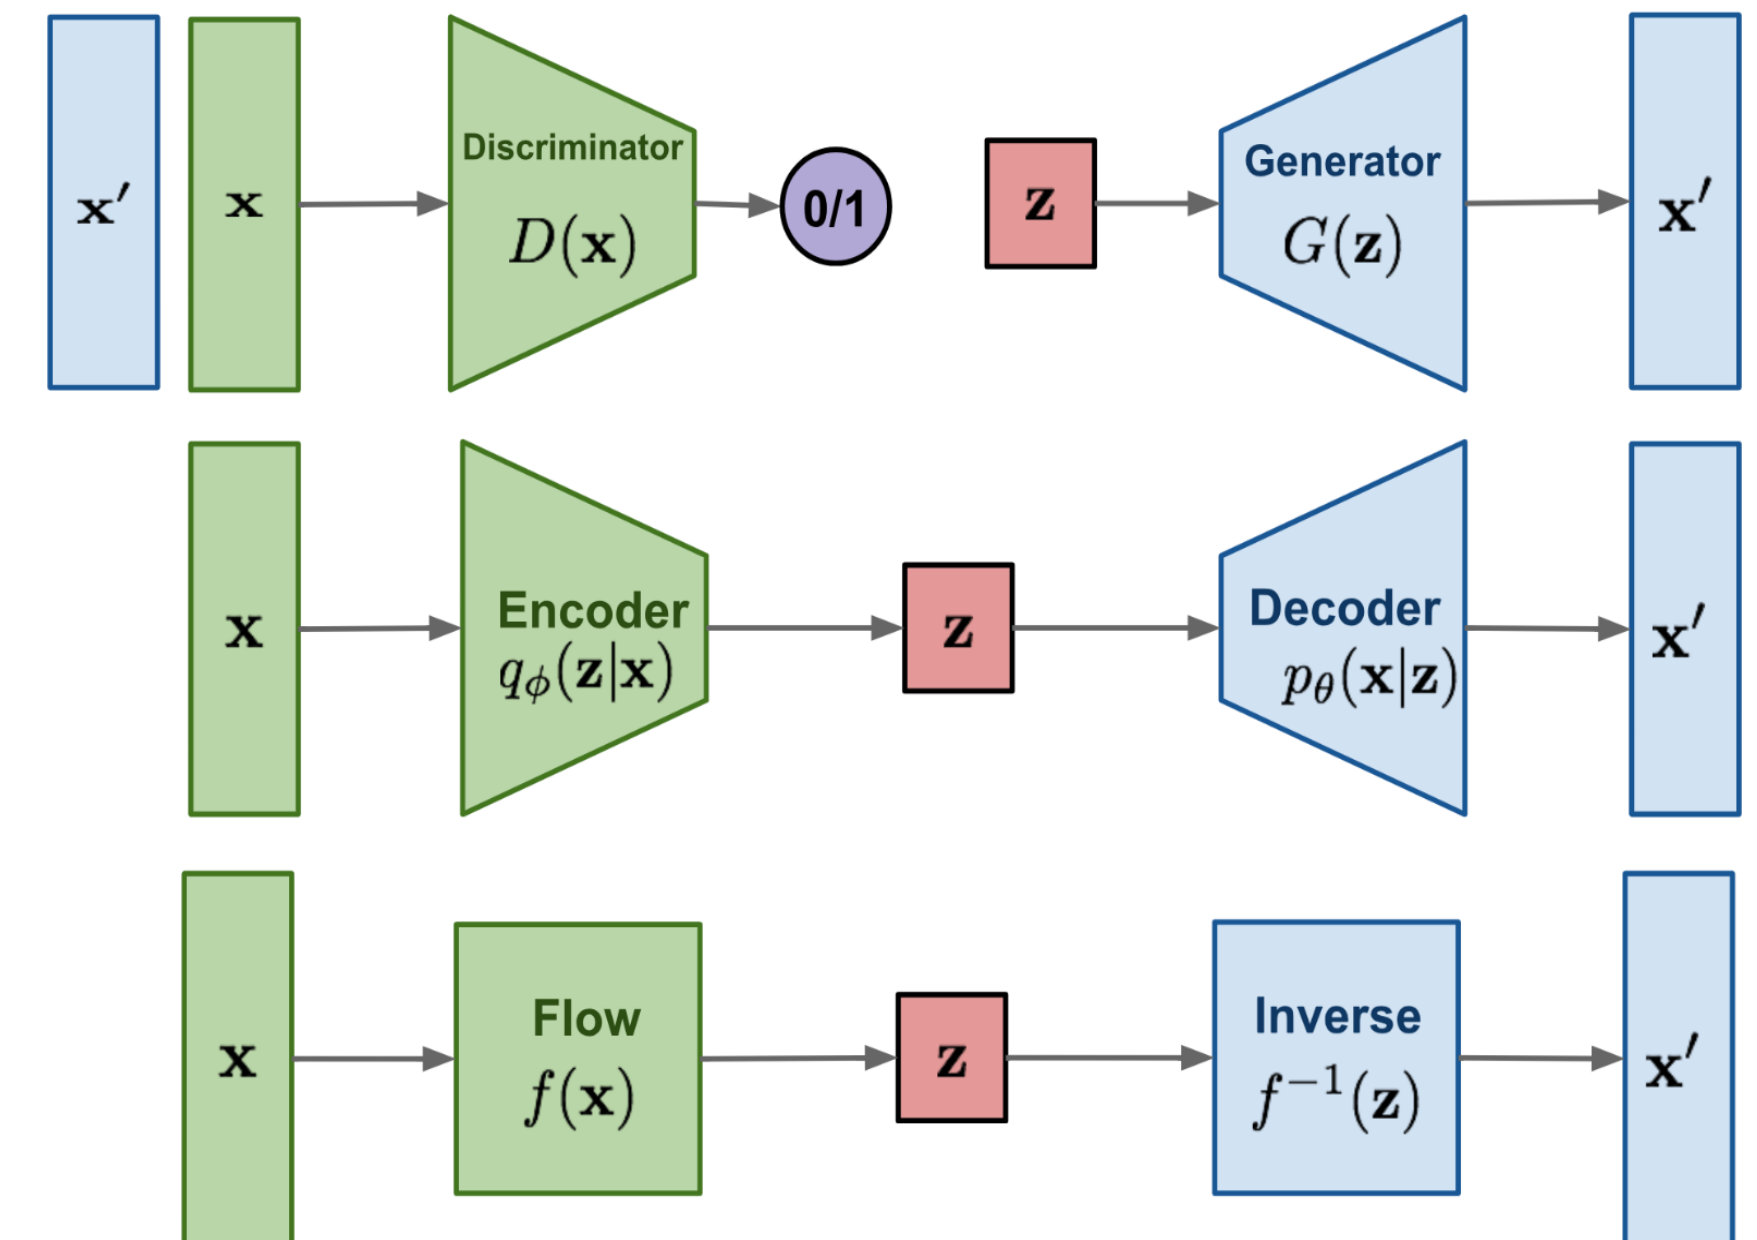
Normalizing Flows (NF)

GANs and VAEs have shown awe-inspiring results for learning complex data distributions and having simple inference methods. However, Neither of them explicitly learns the probability density function of real data $p(\mathbf{x})$ because it is really hard!


GAN: minimax the classification error loss.

VAE: maximize ELBO.

Flow-based generative models: minimize the negative log-likelihood



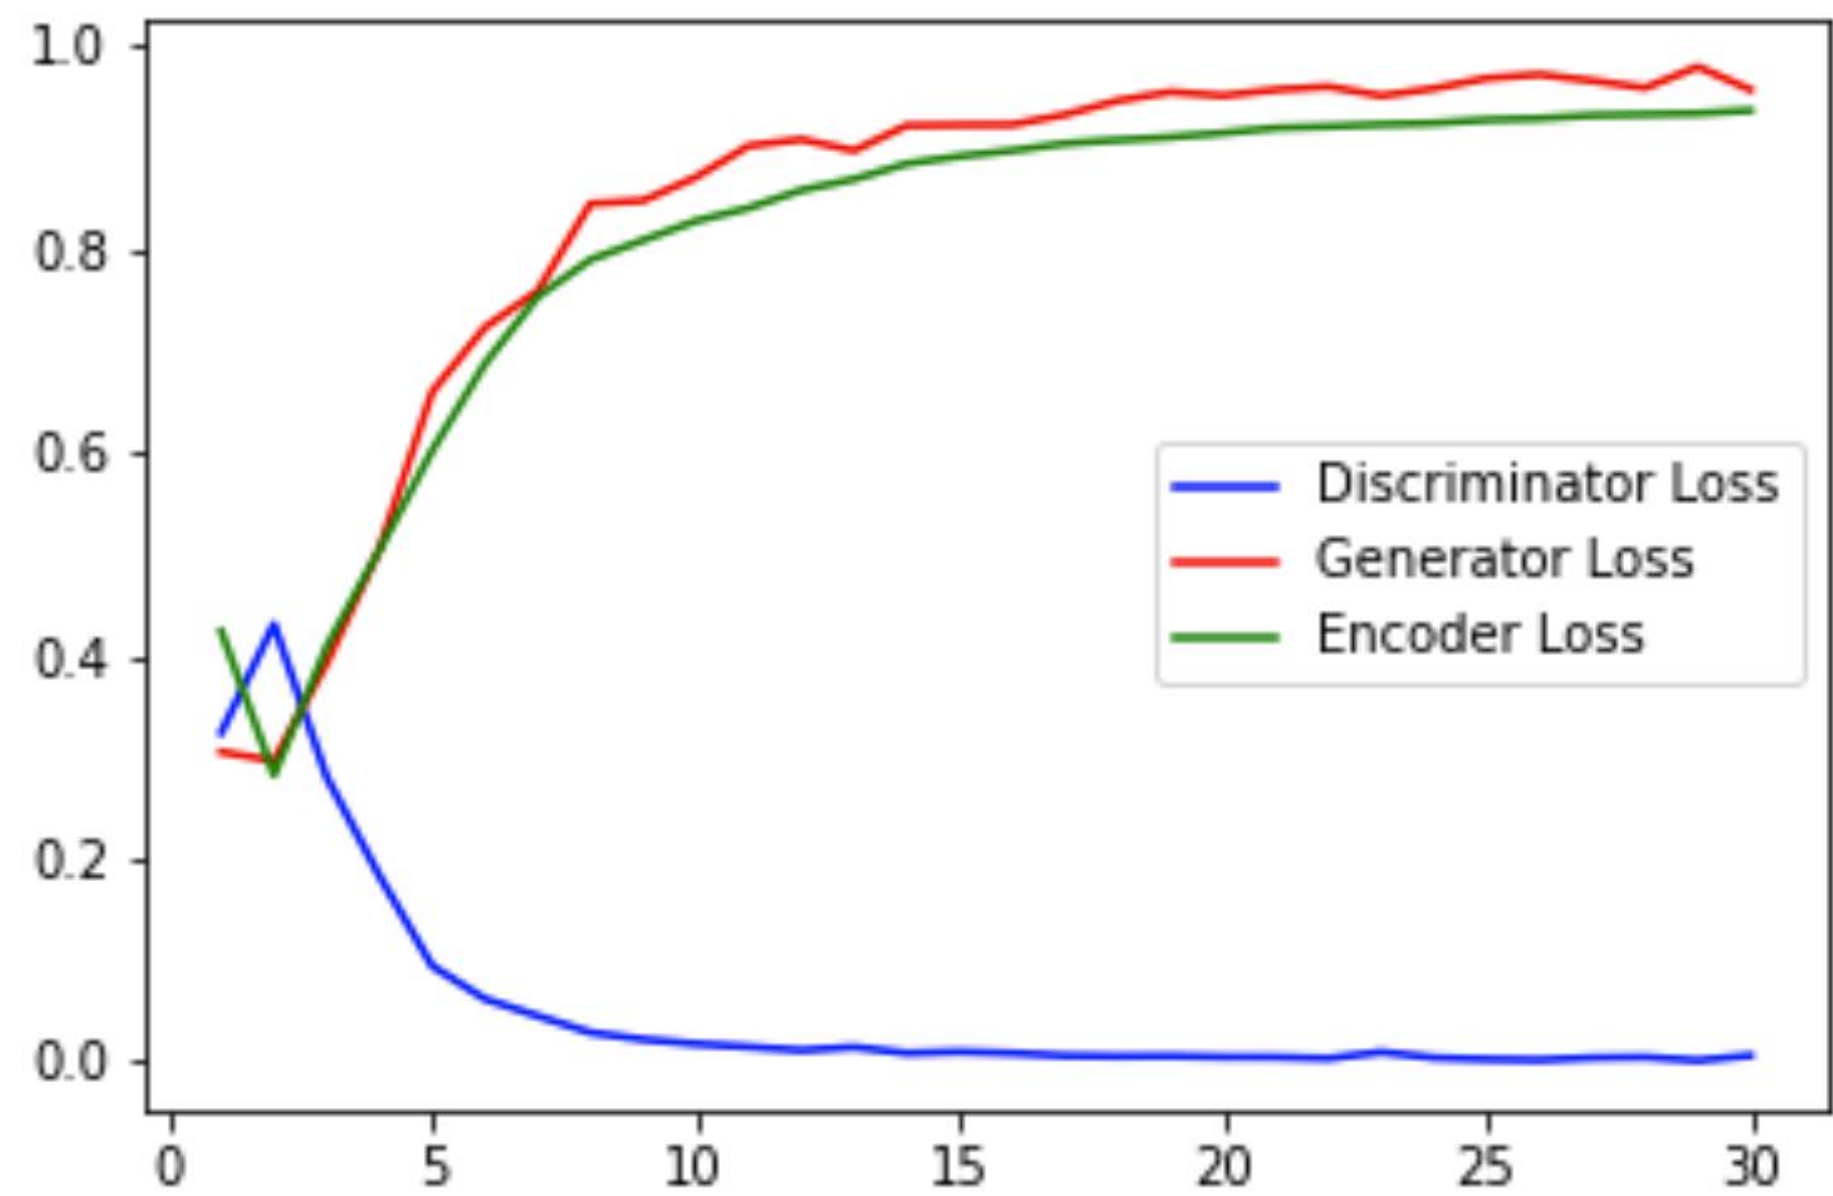
Normalizing Flows (NF)



A normalizing flow transforms a complex distribution into a simple one (Gaussian distribution) by applying a sequence of invertible transformation functions. Flowing through a chain of transformations, eventually we obtain a probability distribution of the final target variable.

Evaluation - BIGAN

BIGAN



	Precision	Recall	F1-Score
Benign	0.514	0.9744	0.673
Attack	0.937	0.293	0.447

USTC-TF16 dataset

Evaluation - Autoencoder

Dataset	Epochs	Train Loss	Validation Loss	Test Loss
CICIDS2017	100	0.01632	0.01631	0.01988
USTC-TF	100	0.01295	0.012977	0.01767

Training, validation and Test losses

	Precision	Recall	F1-Score
Benign	0.8437	0.8635	0.8535
Attack	0.3839	0.347	0.3645

CICIDS-2017 dataset

	Precision	Recall	F1-Score
Benign	0.7654	0.8761	0.817
Attack	0.8008	0.6498	0.7175

USTC-TF16 dataset

Evaluation - Variational Autoencoder

Dataset	Epochs	Train Loss	Validation Loss	Test Loss
CICIDS2017	100	0.03338	0.03345	0.03521
USTC-TF	100	0.03591	0.03583	0.03989

Training, validation and Test losses


	Precision	Recall	F1-Score
Benign	0.8437	0.8635	0.8535
Attack	0.3839	0.347	0.3645

CICIDS-2017 dataset

	Precision	Recall	F1-Score
Benign	0.7506	0.8431	0.7942
Attack	0.7562	0.6347	0.6901

USTC-TF16 dataset


Evaluation- Adversarial Autoencoder



	Precision	Recall	F1-Score
Benign	0.862	0.848	0.855
Attack	0.421	0.45	0.43

CICIDS-2017 dataset

Evaluation – Adversarial Autoencoder



	Precision	Recall	F1-Score
Benign	0.89	0.99	0.938
Attack	0.86	0.84	0.908

USTC-TF16 dataset

Evaluation - Normalizing Flows

	Precision	Recall	F1-Score
Benign	0.8	0.89	0.84
Attack	0.97	0.94	0.96

CICIDS-2017 dataset

	Precision	Recall	F1-Score
Benign	1.00	0.92	0.96
Attack	0.94	1.00	0.97

USTC-TF16 dataset

Training NSF(Neural Spline Flows) on 5% of the datasets:

Conclusion & Perspectives

Short term perspectives:

- Add other datasets and compare the behavior of each model on all datasets.
- Try to generalize our models to make them work for different datasets.

Long term perspectives:

- Transfer the current work to a Federated setting and investigate the collaboration side of FL.

THANK YOU FOR YOUR ATTENTION!

Deep Learning for Anomaly Detection



Meryem Janati Idrissi



Ismail Berrada



JDD'22